

## Module 02

# Talking to AI effectively

Most people underuse AI because they talk to it like a search engine. This module teaches you how to actually communicate with AI, the skill that separates people who get useful results from those who don't.

By Yulia Brusova

18 min read · Reviewed October 2026

### WHAT YOU'LL BE ABLE TO DO

- 1 Write a prompt that includes role, task, context, and format instructions
- 2 Use follow up prompts to refine and improve AI output rather than starting over
- 3 Set up a system prompt or custom instructions for a recurring library task
- 4 Apply a step by step (chain of thought) prompt with explicit criteria to a library evaluation or decision making task
- 5 Identify at least three patron privacy risks in AI prompting and apply strategies to mitigate them

*The most significant skill gap among librarians beginning to use AI is not about knowing which tool to select. It is about knowing how to communicate with it. A vague prompt produces a vague answer; a specific, context rich prompt produces something one can actually use. This module examines the structure of effective prompting, the patterns that serve recurring library tasks, and the practices, including patron privacy, that distinguish responsible AI use from careless use.*

## The anatomy of a good prompt

A useful prompt has four elements. Not every task requires all four, but understanding them allows a librarian to construct prompts deliberately rather than by trial and error.

**Role:** Tell the AI who it is. For example: "You are an experienced academic librarian at a community college with a large first generation student population." Role instructions orient the model toward the vocabulary and concerns of the professional context.

**Task:** Tell it exactly what you want: "Write a 200-word email to faculty explaining our new database access policy." The task should be specific enough that there is only one reasonable interpretation of it.

**Context:** Give it the information it needs: "The instructor has not responded to previous outreach and may be skeptical about library instruction." Context is where most librarians underinvest, and the more relevant background the model has, the less generic the output will be.

**Format:** Tell it how to structure the output: "Friendly but professional. Three short paragraphs. No bullet points." Without format instructions, AI defaults to whatever is typical for the task type.

Consider the difference between "Write me an email about our library databases" and "You are a reference librarian at a community college. Write a friendly 150-word email to first year students introducing three databases they will use for English Composition papers: JSTOR, Academic Search Complete, and ProQuest. Include one sentence about how to get help. No jargon." The second prompt produces a draft that is close to usable, which the librarian then reviews and refines rather than rewriting from scratch. Such specificity is the core skill this module develops.

## A framework for description: product, process, and performance

Dakan and Feller organize communication with AI into three layers (Dakan & Feller, AI Fluency: Framework & Foundations, Anthropic Academy, 2025). The first is **product description**: what exactly should the AI produce, for whom, and for what purpose? For example, asking for a research guide aimed at adult returning students who have never used a library database gives the AI something workable; asking for "a library guide" does not. The second is **process description**: how should the AI approach the task? This might mean asking it to consider three possible structures before choosing one, or to flag uncertainty rather than fill a gap silently. The third, and the one most often skipped, is **performance description**: how should the AI behave throughout the exchange? Should it challenge the framing or follow the lead, answer briefly or in detail, show its reasoning or only the result? Defining all three layers before a complex task takes a minute and reliably saves considerably more.

## Iteration is the skill

Most users send one prompt, receive one response, and either accept it or abandon the attempt. Iteration, the act of following up, refining, and redirecting, is where the genuine value lies. Useful follow ups include: "Make this two paragraphs instead of four"; "This sounds too formal: make it warmer"; "The second section does not address what I need; here is what I require: [specifics]"; and "Give me three alternative opening sentences." One follow up worth making a habit is "Identify any claims in this draft that are unsupported and tell me what needs verification."

In order to iterate effectively, one should continue the existing conversation rather than starting a new one, because the AI retains context within a session. For example, once a librarian has established that the audience is returning adult students who are skeptical of library instruction, every follow up benefits from that context without the librarian repeating it.

## Prompt patterns for common library task types

The most efficient way to develop prompting skill is to work from proven structures that can be adapted to recurring tasks rather than written from scratch each time. Several patterns serve library practice reliably:

**The draft-and-refine pattern** (emails, instruction content, policy text): "You are a [role]. Write a [format] for [audience] that [purpose]. Tone: [description]. Length: [target]. Include: [specific elements]. Exclude: [things to omit]." This is the workhorse pattern for most library communication.

**The summarize-and-extract pattern** (long documents): "Here is a [document type]. Summarize it in [number] sentences for [audience]. Then extract: (1) the key policy changes, (2) any deadlines, (3) any action required from library staff. Flag anything that is unclear." This pattern is especially useful for accreditation reports, vendor contracts, and administrative memos.

**The explain-to-audience pattern** (instruction and patron communication): "Explain [concept] to a [audience description] who [relevant context about their knowledge level]. Use plain language. No jargon. Give one concrete example. Limit to [word count]."

**The generate-options pattern** (brainstorming): "Generate [number] different [options] for [goal]. Each option should be distinct. Format as a numbered list with a one sentence rationale for each." For example, five research question framings for a patron who is stuck, or ten titles for a workshop series.

Patterns that produce reliable results are worth saving in a shared document accessible to the whole library team. Such a collection serves as an institutional knowledge base that outlasts any individual's tenure, and it is the foundation for the prompt library covered in Module 10.

## Chain of thought prompting for complex decisions

Asking the AI to reason through a problem step by step, against explicit criteria, before producing a conclusion often improves output on complex tasks. For example, instead of

asking "Should we purchase a subscription to this new database?", a more effective prompt is: "I am evaluating a new database subscription for our academic library. Here is the vendor's pitch and pricing: [paste]. Work through the following considerations step by step, and set out your findings, uncertainties, and open questions before giving a recommendation: (1) alignment with our subject areas, (2) cost relative to comparable resources, (3) data privacy terms, (4) ILL implications, (5) likely faculty and student usage." The response is usually more structured, more complete, and easier to act on than an answer to the shorter question, and it lets the librarian see whether the AI has understood the question and where information is missing.

A detailed explanation is not proof of accuracy, however. The stated reasoning can sound convincing while resting on a wrong fact, and it does not necessarily reflect how the model actually reached its answer, so important claims still require verification. Additionally, many current models include a built-in reasoning or extended thinking mode that does some of this work automatically, which makes the explicit criteria in the prompt more valuable than the phrase "think step by step" itself.

## **System prompts and custom instructions**

Most AI tools allow users to set persistent instructions that apply to every conversation automatically, which removes the need to re-explain one's context in each session. Useful instructions for library work include institution type and student population ("I work at a community college serving many first generation students and students returning to higher education"), professional role ("I am a reference and instruction librarian"), preferred output style ("Use plain language and active voice; avoid jargon"), and explicit expectations ("Distinguish verified facts from suggestions, and flag claims that need an authoritative source").

In ChatGPT, this feature appears as Custom Instructions in the personalization settings. In Claude, it takes the form of profile preferences, project instructions inside a Project, and styles. In Gemini, the equivalents are saved info and Gems, reusable custom versions of Gemini with their own instructions. Claude's Projects also allow the librarian to attach documents, such as a collection development policy or a list of available databases, that persist across every conversation in the project. Feature names and availability change often and vary by plan, so it is worth checking the current settings. Furthermore, persistent context improves consistency but does not remove the need to check each response, and the instructions themselves should be reviewed periodically as roles, resources, and policies change.

## When to give AI information versus ask for information

There are two fundamentally different modes of prompting, and knowing which one a task requires changes how one writes the prompt and evaluates the output.

**Information in:** The librarian shares a document or draft and asks AI to work with it. For example: "Here is a LibGuide I drafted. Improve the clarity of the introduction and suggest more descriptive section headings." This mode is considerably more reliable, because every statement in the output can be checked against the source. It does not eliminate risk: a model can still misread a document, omit an important detail, or add a claim the source does not support.

**Information out:** The librarian asks AI to produce information it does not have in front of it. For example: "What are the most useful databases for nursing research?" This mode requires more verification, because the model draws on training data that may be outdated or imprecise about specialized library resources, and it may confidently name databases that have changed, merged, or been discontinued.

For drafting, editing, summarizing, and brainstorming, the default should be the information in mode: paste the document, share a draft rather than asking for one from scratch, and provide the patron's question rather than a paraphrase, with identifying details removed. Such an approach produces better output and makes verification considerably easier, and it reflects a sound principle: AI is most useful as a collaborator working with what the librarian brings, not as an independent source of facts.

## Common prompting mistakes and how to avoid them

Certain prompting mistakes are nearly universal among librarians beginning to use AI, and naming them directly saves the frustration of discovering them one by one.

**Prompts that are too short.** A two or three word query such as "email about databases," the kind of thing one types into Google, produces a generic response. A useful prompt is almost always several sentences long, with role, task, context, and format specified.

**Accepting the first response.** The first response is a starting point, not the output. Targeted follow ups, as described above, are where the quality gain happens.

**Asking for too many things at once.** A prompt that asks for an instruction email, a lesson plan, and five assessment activities will produce mediocre results across all three. Asking for

one thing, refining it, and then moving on also allows the librarian to verify each stage before the next.

**Re-explaining context in every new conversation.** Custom instructions, described above, remove this inefficiency, and they remain one of the most underused features among librarians who use AI regularly.

Several smaller fixes are worth knowing (Dakan & Feller, Anthropic Academy, 2025). When a response is the wrong length, a concrete target such as "two paragraphs" or "under 100 words" works better than "make it shorter." When the tone is wrong, showing works better than describing: pasting an example of the library's existing voice and asking the AI to match it is more reliable than adjectives like "more conversational." And when an interaction is not producing what is needed, the most efficient move is often to describe the actual goal and ask the AI to help rewrite the prompt itself.

## Evaluating AI for your specific workflows

AI is not uniformly capable across tasks. A model that drafts excellent reference emails may struggle with subject specific instruction content or specialized policy questions, and the only way to know how it performs on a given workflow is to test it directly.

A practical approach is to choose one recurring task, gather five to ten examples of it, run them through the AI with consistent instructions, and compare the results with one's own work, noting errors, omissions, the corrections required, and the time saved. For example, a cataloging librarian who tests AI generated subject heading suggestions against records they have already cataloged develops firsthand knowledge of its accuracy in their own collection areas. Such testing requires no technical infrastructure, only the discipline to compare AI output against professional judgment on tasks one knows well. The result is not a pass or fail verdict on the tool but a calibrated sense of where AI assistance is worth incorporating, and where the verification it requires would cancel out the time it saves.

## Patron privacy and what not to share in a prompt

Libraries have a long and principled commitment to patron privacy, and that commitment extends to AI use. If a reference librarian pastes a patron's question, along with identifying details about the patron's research, into a commercial AI tool, that information may be processed on the company's servers, logged, and potentially used for training or reviewed by employees, depending on the tool's terms and settings. Such a practice would breach patron privacy in most library contexts, particularly where patron records carry legal protection.

As a general rule, remove all identifying information before pasting patron questions into AI, including indirect details, such as a rare course or a distinctive research topic, that could identify a patron in combination. For example, instead of submitting "A nursing student named Maria asked about finding articles on postpartum depression for her capstone project at St. Louis Community College," the prompt should read: "A nursing student is looking for peer reviewed articles on postpartum depression for a capstone project. Suggest three search strategies and two relevant databases." Anonymization alone does not make every use permissible, however: confidential records and sensitive research data should not go into an unapproved tool at all, and when the requirements are unclear, the right step is to consult the appropriate staff first.

Furthermore, librarians should understand the privacy terms of any AI tool used for professional work. Free consumer tiers typically reserve more rights over user inputs than institutional plans do. ChatGPT's Business and Enterprise plans (Business was called Team until August 2025), Claude's Team and Enterprise plans, and Google Workspace editions of Gemini all offer stronger data protection terms, including commitments not to use inputs for training. The conversation with one's institution about which AI tools are approved for professional use is as important as any prompting skill, and librarians are well positioned to lead it.

#### **PRACTITIONER NOTE**

A reliable prompt practice for reference work involving unclear patron questions: when a patron's research need is underspecified, remove any identifying details and paste the question into Claude with the following prompt: 'A patron has asked the following research question. Generate five clarifying questions I could ask to better understand their actual need: [question].' This approach surfaces angles that would not have been considered independently, particularly for research topics outside one's subject expertise. Such a practice takes approximately thirty seconds and consistently improves the quality of the reference interaction that follows, which is a reasonable return on a very small investment of time.

#### **KEY TAKEAWAYS**

- Effective prompts include four elements: role, task, context, and format. Specificity is the core skill.

- Iteration through follow up prompts produces better results than accepting a first response or starting over.
- System prompts and custom instructions eliminate the need to restate your professional context in every session.
- Default to the information in mode: give AI your documents to work with, then check the output against them.
- Patron privacy obligations apply to prompts: strip direct and indirect identifiers, and keep confidential records out of unapproved tools.
- Chain of thought prompting against explicit criteria makes complex decisions easier to review, though a detailed explanation is not proof of accuracy.

---

## References

### APA 7TH EDITION

1. Dakan, R., & Feller, A. (2025). *AI fluency: Framework & foundations* [Online course]. Anthropic Academy. <https://anthropic.skilljar.com/ai-fluency-framework-foundations>

### ACRL AI Competencies covered

Knowledge & Understanding

Use & Application

Sub-competencies: 4.3, 2.1 · ACRL AI Competencies (2025)